

Project 3: Data Manipulation with Pandas

John Wesley Mathis

Dr. Anthony Choi

June 08, 2024

ECE/SSE 591, Summer 2024

Table of Contents:

Deliverable Table.....	5
1. Introduction.....	6
2. Top Movie Data Analysis.....	7
3. Conclusion.....	33
References.....	34

Table of Figures

Figure 1: Using Pandas objects and importing CSV files.....	7
Figure 2: Collecting information about the data.....	7
Figure 3: Collecting information about the dataset.....	8
Figure 4: Changing the index of the dataset.....	9
Figure 5: Adding a Rank column.....	10
Figure 6: Moving 'Rank' column to the beginning of the dataset.....	11
Figure 7: Sorting the data by IMDB rating.....	12
Figure 8: Finding what data is missing.....	13
Figure 9: Removing unused columns from the dataset.....	14
Figure 10: Dealing with missing data.....	15
Figure 11: Checking for duplicated data.....	16
Figure 12: Converting 'Gross' column from String to Integer.....	16
Figure 13: Cleaning up the 'Runtime' column.....	17
Figure 14: Converting release year to datetime format.....	18
Figure 15: Cleaning up the genre column.....	18
Figure 16: Determining top movie genres.....	19
Figure 17: Determining top 5 most voted movies.....	19
Figure 18: Determining bottom 5 least voted movies.....	20
Figure 19: Determining which movie has the longest runtime.....	20
Figure 20: Determining which movie has the shortest runtime.....	21
Figure 21: Using `query()` to filter data.....	21
Figure 22: Using `eval()` to calculate new data.....	22
Figure 23: Determining which year had the most top movies.....	22
Figure 24: Determining highest grossing movie.....	23
Figure 25: Determining which genre has the IMDB rating and which has the highest gross.....	24
Figure 26: Discovering which directors consistently produce top movies using pivot tables.....	25

Figure 27: Demonstrating hierarchical indexing.....	26
Figure 28: Importing two new datasets and showing information about dataset 1.....	27
Figure 29: Dataset 2 information.....	27
Figure 30: Merging the two datasets on 'id'.....	28
Figure 31: Removing unused columns.....	29
Figure 32: Setting the index to 'genres' and 'release_date'	30
Figure 33: Code for vectorized string operations.....	30
Figure 34: Output to Figure 32 code.....	31
Figure 35: Demonstrating `concat()`	32
Figure 36: Using `concat()` to append new data.....	33

Deliverable Table

The purpose of this table is to provide a complete view of the concepts covered in chapter 3 of *"Python Data Science Handbook"* (VanderPlas, 2016) and provide a general page location for where the topic was demonstrated.

Deliverables	Location
Introducing Pandas Objects	6
Data Indexing and Selection	6
Operating on Data in Pandas	6-7
Handling Missing Data	12-14
Hierarchical Indexing	24
Combining Datasets: Concat and Append	31-32
Combining Datasets: Merge and Join	25-30
Aggregation and Grouping	22
Pivot Tables	23
Vectorized String Operations	25-30
Working with Time Series	16
High-Performance Pandas: eval() and query()	20-21

Additionally, here is a link to my GitHub where the datasets and the Jupyter Notebook for the project can be downloaded: https://github.com/jwmathis/SSE591_Project3.git

1. Introduction

Python has a rich repository of libraries that aid scientists and researchers in data analysis and manipulation. One of the most common libraries in use is Pandas, which is built on top of NumPy and provides a higher-level, and more flexible interface for data handling. While NumPy excels at efficient numerical computations with arrays, Pandas introduces data structures like Series and DataFrame that offer a more intuitive means to work with structured data.

Because of Pandas' Series and DataFrame objects, data scientists have an indispensable tool to handle, clean and manipulate data in tabular form. These objects support a wide range of operations, from simple data aggregation and filtering to complex time-series analysis. The library's ability to handle missing data, merge datasets, and perform group-by operations adds significant value to Python's data manipulation kit.

This report aims to demonstrate my proficiency in Python data manipulation techniques as covered in Chapter 3 of the "Python Data Science Handbook" by Jake VanderPlas (2016). This report attempts to illustrate the core concepts and functionalities of the Pandas library by implementing the concepts into a single project. The code presented in this report was developed using Visual Studio Code with Jupyter Notebook extensions. I will provide detailed explanations, highlighting key features and operations that make Pandas an essential tool for data analysis.

2. Top Movie Data Analysis

I chose the IMDB Top 1000 movies dataset to analyze. This dataset contains information about the top 1000 movies. After importing the dataset, I viewed the first few rows by outputting it to the screen using `'head()'`. Figure 1 below shows the code and output.

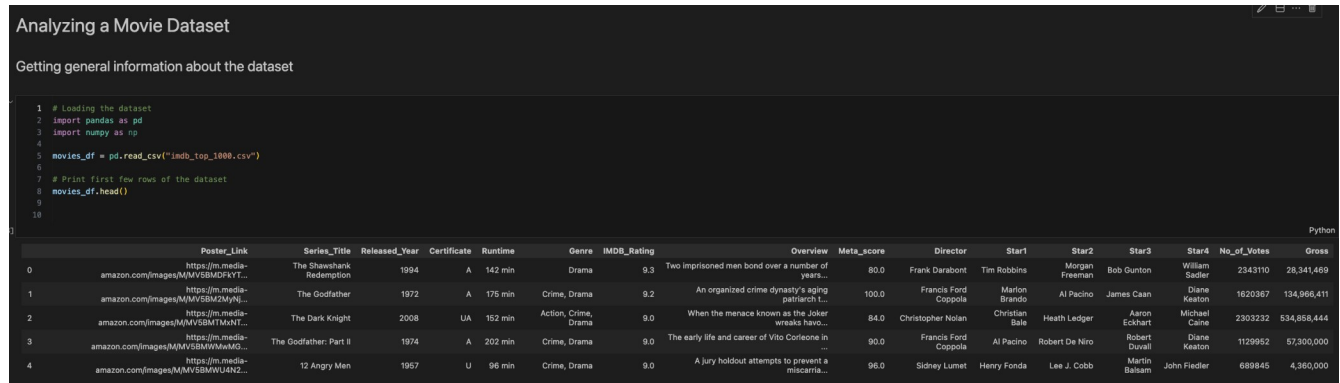


Figure 1: Using Pandas objects and importing CSV files

Using `shape()`, `size()`, and `describe()` I obtained summary statistics of the numerical columns. There are only three numerical columns, IMDB rating, metascore, and number of votes. For this particular data set this doesn't tell me to much. However, we find out that the average IMDB rating is a 7.94 and the average metascore for movies is a 77.97. However, because we don't know much about the data, we don't understand how skewed this information may be. Figure 2 shows the code and output.

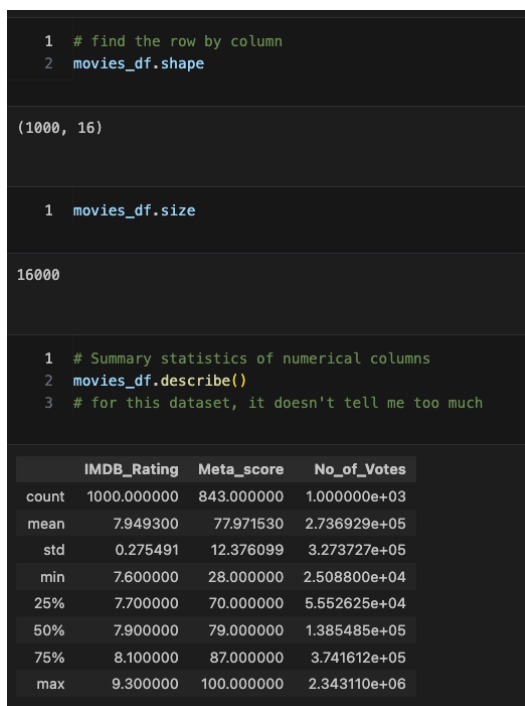


Figure 2:

Collecting information about the data

To better understand the data, I used `info()` to obtain general information about the DataFrame. From this view, I learned the column names, the data type for each column, and if each column contained all the information. The *Certificate*, *Metascore*, and *Gross* columns were missing some data. Figure 3 shows the code and output.

```
1 # Get general information about the dataframe
2 movies_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Poster_Link     1000 non-null   object
1   Series_Title    1000 non-null   object
2   Released_Year   1000 non-null   object
3   Certificate      899 non-null    object
4   Runtime         1000 non-null   object
5   Genre           1000 non-null   object
6   IMDB_Rating     1000 non-null   float64
7   Overview        1000 non-null   object
8   Meta_score      843 non-null    float64
9   Director        1000 non-null   object
10  Star1            1000 non-null   object
11  Star2            1000 non-null   object
12  Star3            1000 non-null   object
13  Star4            1000 non-null   object
14  No_of_Votes     1000 non-null   int64
15  Gross            831 non-null    object
dtypes: float64(2), int64(1), object(13)
memory usage: 125.1+ KB
```

Figure 3: Collecting information about the dataset

Now with this understanding of the data, I begin to clean it up. Because the data seems to be un-ranked, just simply the top 1000 movies in an arbitrary order, I changed the index from the basic 0-999 using the `set_index()` method. I changed the index to be the series title. Figure 4 shows this output.

```
1 movies_df.set_index('Series_Title', inplace=True) # Change index to be Title column
2 movies_df.head(4)
```

Series_Title	Poster_Link	Released_Year	Certificate	Runtime	Genre	IMDB_R
The Shawshank Redemption	https://m.media-amazon.com/images/M/MV5BMDFKYT...	1994	A	142 min	Drama	
The Godfather	https://m.media-amazon.com/images/M/MV5BM2MyNj...	1972	A	175 min	Crime, Drama	
The Dark Knight	https://m.media-amazon.com/images/M/MV5BMTxNT...	2008	UA	152 min	Action, Crime, Drama	
The Godfather: Part II	https://m.media-amazon.com/images/M/MV5BMWwMG...	1974	A	202 min	Crime, Drama	

Figure 4: Changing the index of the dataset

After viewing the first few rows of data and the last few rows using the `head()` and `tail()` methods respectively, I confirmed that the data was un-ranked. The dataset is simply a list of the top 1000 movies. After reviewing the data info from before (Fig 4), I decided to add a rank column as seen in Figure 5.

```
Cleaning up the dataset

• adding columns before analysis
• removing rows
• cleaning up strings
• converting data types

1 # Add a 'Rank' column
2 movies_df['Rank'] = range(1, len(movies_df) + 1)
3 movies_df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 1000 entries, The Shawshank Redemption to The 39 Steps
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Poster_Link           1000 non-null   object
1   Released_Year         1000 non-null   object
2   Certificate           899 non-null    object
3   Runtime               1000 non-null   object
4   Genre                 1000 non-null   object
5   IMDB_Rating           1000 non-null   float64
6   Overview              1000 non-null   object
7   Meta_score            843 non-null    float64
8   Director              1000 non-null   object
9   Star1                 1000 non-null   object
10  Star2                 1000 non-null   object
11  Star3                 1000 non-null   object
12  Star4                 1000 non-null   object
13  No_of_Votes           1000 non-null   int64
14  Gross                 831 non-null    object
15  Rank                  1000 non-null   int64
dtypes: float64(2), int64(2), object(12)
memory usage: 132.8+ KB
```

Figure 5: Adding a Rank column

However, there are a lot of columns in this dataset, and I wanted to be able to quickly and easily view the ranks. So I retrieved the columns and converted them to a list and stored the information in a variable for access. I removed the string 'Rank' from the list and inserted it again into the list at position 0. Then I re-ordered the DataFrame using the variable that stored the columns. Figure 6 shows the output.

```
1 # get list of column names
2 new_columns = list(movies_df.columns)
3 # remove the 'Rank' column from the list
4 new_columns.remove('Rank')
5 # insert 'Rank' column at specified index
6 new_columns.insert(0, 'Rank')
7 # Reorder the DataFrame Columns
8 movies_df = movies_df[new_columns]
9 # Check DataFrame info
10 movies_df.info()
```

<class 'pandas.core.frame.DataFrame'>
Index: 1000 entries, The Shawshank Redemption to The 39 Steps
Data columns (total 16 columns):
Column Non-Null Count Dtype
--- ---
0 Rank 1000 non-null int64
1 Poster_Link 1000 non-null object
2 Released_Year 1000 non-null object
3 Certificate 899 non-null object
4 Runtime 1000 non-null object
5 Genre 1000 non-null object
6 IMDB_Rating 1000 non-null float64
7 Overview 1000 non-null object
8 Meta_score 843 non-null float64
9 Director 1000 non-null object
10 Star1 1000 non-null object
11 Star2 1000 non-null object
12 Star3 1000 non-null object
13 Star4 1000 non-null object
14 No_of_Votes 1000 non-null int64
15 Gross 831 non-null object
dtypes: float64(2), int64(2), object(12)
memory usage: 132.8+ KB

Figure 6: Moving 'Rank' column to the beginning of the dataset

After re-ordering the DataFrame, I used `sort_values()` using the IMDB rating to sort the movies in ascending order. Then I set the index to be the *Rank* column. Figure 7 shows the code and output

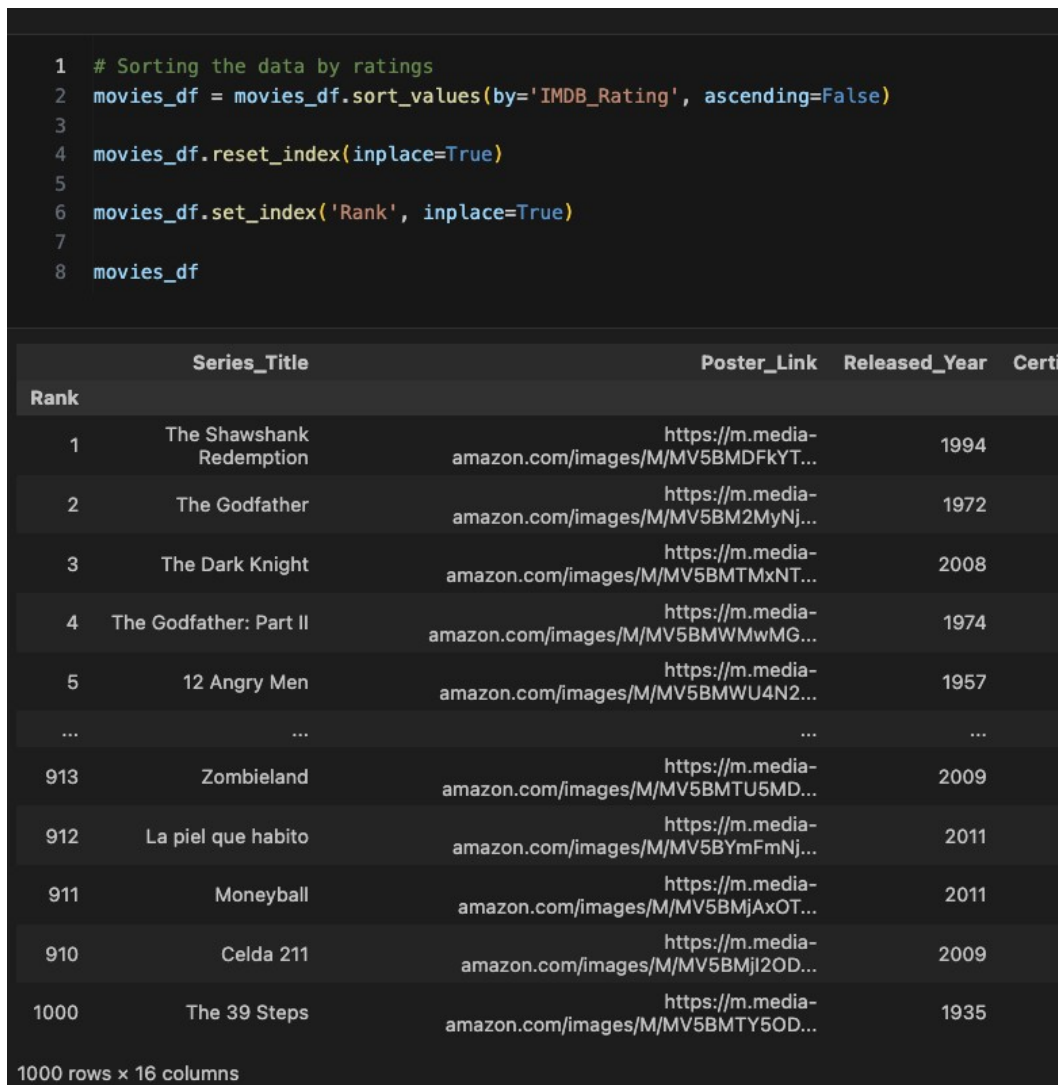


Figure 7: Sorting the data by IMDB rating

From figure ## I already saw that some of the data was missing. So I wanted to clean that up some. So I first checked what columns contain null values using `isna()` and `sum()` to get a number of missing data in each column. From this I verify that the only columns missing data are *Certificate*, *Meta_score*, and *Gross*. Figure 8 shows the code and results.

```
1 # Handling missing data
2 print(movies_df.isnull().sum())
```

✓ 0.0s

Series_Title	0
Poster_Link	0
Released_Year	0
Certificate	101
Runtime	0
Genre	0
IMDB_Rating	0
Overview	0
Meta_score	157
Director	0
Star1	0
Star2	0
Star3	0
Star4	0
No_of_Votes	0
Gross	169
dtype:	int64

Figure 8: Finding what data is missing

I didn't want the Certificate column or the Poster_Link column, so I decided to remove both. Figure 9 shows the code and the output. I remove these columns by using `drop()`. With *Certificate* gone, there is one less column of missing data to deal with.

```
1 # drop unused columns
2 movies_df.drop(['Certificate', 'Poster_Link'], axis=1, inplace=True)
3 movies_df.info()

✓ 0.0s

<class 'pandas.core.frame.DataFrame'>
Index: 1000 entries, 1 to 1000
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Series_Title     1000 non-null   object
1   Released_Year    1000 non-null   object
2   Runtime          1000 non-null   object
3   Genre            1000 non-null   object
4   IMDB_Rating      1000 non-null   float64
5   Overview         1000 non-null   object
6   Meta_score       843 non-null    float64
7   Director         1000 non-null   object
8   Star1            1000 non-null   object
9   Star2            1000 non-null   object
10  Star3            1000 non-null   object
11  Star4            1000 non-null   object
12  No_of_Votes      1000 non-null   int64
13  Gross            831 non-null    object
dtypes: float64(2), int64(1), object(11)
memory usage: 117.2+ KB
```

Figure 9: Removing unused columns from the dataset

Gross represents the amount of money the movie made. Rather than finding appropriate data to fill in these missing values for the movies, I decide to simply remove all rows that are missing data in either the metascore column or the gross column. I used `dropna()` to remove these rows. Figure 10 shows the code and output. This reduces the number of entries to 750, meaning I lost 25% of my data.

```
1 #movies_df.fillna(0)
2 #print(movies_df.isna().sum())
3 #lets remove rows with null data
4 movies_df.dropna(inplace=True)
5 movies_df.info()

✓ 0.0s

<class 'pandas.core.frame.DataFrame'>
Index: 750 entries, 1 to 911
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Series_Title          750 non-null    object
1   Released_Year         750 non-null    object
2   Runtime               750 non-null    object
3   Genre                 750 non-null    object
4   IMDB_Rating           750 non-null    float64
5   Overview              750 non-null    object
6   Meta_score            750 non-null    float64
7   Director              750 non-null    object
8   Star1                 750 non-null    object
9   Star2                 750 non-null    object
10  Star3                 750 non-null    object
11  Star4                 750 non-null    object
12  No_of_Votes           750 non-null    int64
13  Gross                 750 non-null    object
dtypes: float64(2), int64(1), object(11)
memory usage: 87.9+ KB
```

Figure 10: Dealing with missing data

I wanted to also check if there are any duplicates in the dataset. Figure 11 shows the code and output. I used `'duplicated()'` to determine this. The result return 0, meaning there are no duplicated rows.

```

1 # determine if there is any duplicated data
2 sum(movies_df.duplicated())
✓ 0.0s
0

```

Figure 11: Checking for duplicated data

Next, I attempted to convert the data in Gross to numerical data since its data type is still an object. However, this produced a *ValueError*. So I updated the code to coerce errors to NaN. Unfortunately this removed all the data from that column. So I backtracked and decided to clean the strings up. After I cleaned the strings, I then converted the values to numerical values. Figure 12 shows my updated code and output and shows that the Gross column is now a type int64.

```

1 # clean up the strings in Gross column for converting to numerical values
2 movies_df['Gross'] = movies_df['Gross'].replace(r'[\$,]', '', regex=True).str.strip()
3 display(movies_df.head(3)) # check that Gross column has changed
4 movies_df['Gross'] = pd.to_numeric(movies_df['Gross'])
5 movies_df.info()
6
✓ 0.0s

```

Rank	Series_Title	Released_Year	Runtime	Genre	IMDB_Rating	
1	The Shawshank Redemption	1994	142 min	Drama	9.3	Two imprisoned
2	The Godfather	1972	175 min	Crime, Drama	9.2	An organiz
3	The Dark Knight	2008	152 min	Action, Crime, Drama	9.0	When the men

```

<class 'pandas.core.frame.DataFrame'>
Index: 750 entries, 1 to 911
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Series_Title    750 non-null   object
1   Released_Year   750 non-null   object
2   Runtime         750 non-null   object
3   Genre          750 non-null   object
4   IMDB_Rating     750 non-null   float64
5   Overview       750 non-null   object
6   Meta_score     750 non-null   float64
7   Director       750 non-null   object
8   Star1          750 non-null   object
9   Star2          750 non-null   object
10  Star3          750 non-null   object
11  Star4          750 non-null   object
12  No_of_Votes    750 non-null   int64
13  Gross          750 non-null   int64
dtypes: float64(2), int64(2), object(10)
memory usage: 87.9+ KB

```

Figure 12: Converting 'Gross' column from String to Integer

Next, I cleaned up the Runtime column. I wanted it to be numerical data as well. To accomplish this I first removed the word minute from the columns by using `str.replace()`. Then I used `pd.to_numeric()` to convert the strings to numerical values. Figure 13 shows the results.

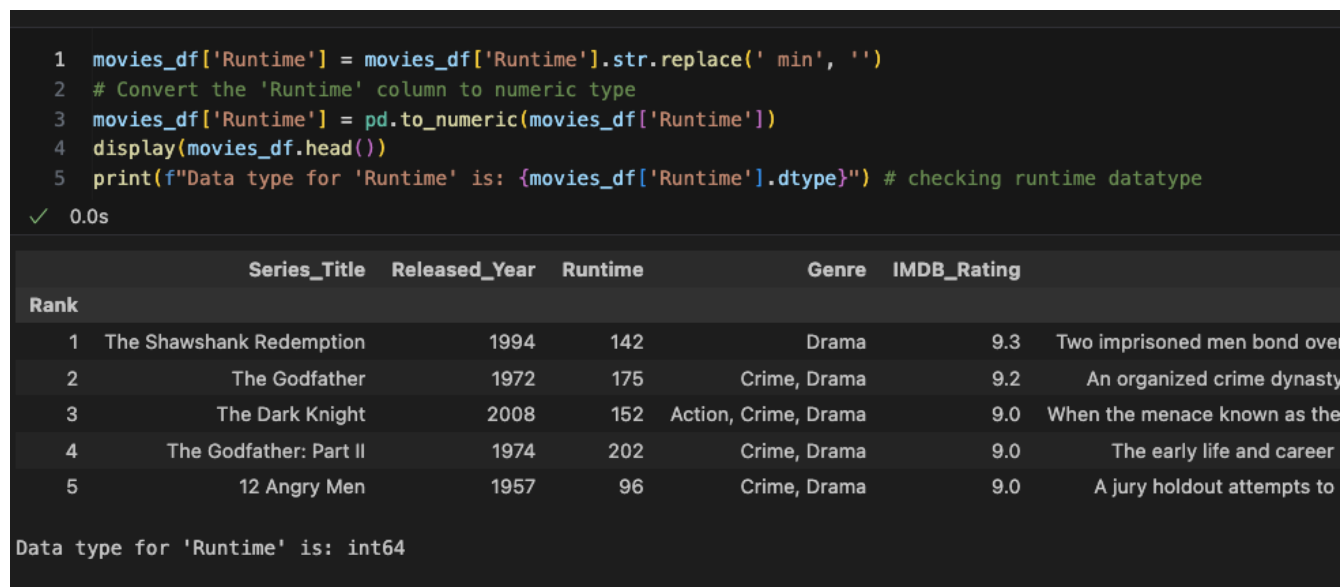


Figure 13: Cleaning up the 'Runtime' column

There is a release year column so I wanted to turn this into the datetime format. Using `pd.to_datetime()` I attempted to do this. Unfortunately, this only works if it is a full date and this column of data only contains the year. So to demonstrate this functionality, I decided to add the full date to the column by filling it with the correct year and a placeholder date of '01-01'. However, all of the rows do not contain a valid year, but it does contain some type of information that is not considered NaN or null. To remove invalid rows, I created a mask to find only the rows with a valid year. Then, I updated the DataFrame to only include these valid year rows. Afterwards, I converted the release year to the proper datetime format. After this was accomplished, I reset the column to show only the year values by using the `dt.year`. Figure 14 shows the code and result.

```

1 # Create a mask to identify rows with valid years (to remove rows with invalid year values)
2 valid_year_mask = movies_df['Released_Year'].str.match(r'^\d{4}$')
3
4 # Filter out rows with invalid years
5 movies_df = movies_df.loc[valid_year_mask].copy()
6
7 # Convert 'Released_Year' to datetime
8 movies_df['Released_Year'] = pd.to_datetime(movies_df['Released_Year'] + '-01-01')
9
10 # Print the updated DataFrame
11 display(movies_df.head(2))
12
13 # Replace dates with just the year value
14 movies_df['Released_Year'] = movies_df['Released_Year'].dt.year
15 display(movies_df.head(2))
16
✓ 0.0s

```

	Series_Title	Released_Year	Runtime	Genre	IMDB_Rating	Overview	Meta_sco
Rank							
1	The Shawshank Redemption	1994-01-01	142	Drama	9.3	Two imprisoned men bond over a number of years...	80
2	The Godfather	1972-01-01	175	Crime, Drama	9.2	An organized crime dynasty's aging patriarch t...	100

	Series_Title	Released_Year	Runtime	Genre	IMDB_Rating	Overview	Meta_sco
Rank							
1	The Shawshank Redemption	1994	142	Drama	9.3	Two imprisoned men bond over a number of years...	80
2	The Godfather	1972	175	Crime, Drama	9.2	An organized crime dynasty's aging patriarch t...	100

Figure 14: Converting release year to datetime format

Lastly, I clean up the genre column by using `str.strip()` and `str.lower()` to ensure there was nothing weird with the column. Figure 15 shows the code and the results.

```

1 # String cleaning
2 movies_df['Genre'] = movies_df['Genre'].str.strip().str.lower()
3 print(movies_df['Genre'].head())

```

✓ 0.0s

```

Rank
1      drama
2  crime, drama
3  action, crime, drama
4  crime, drama
5  crime, drama
Name: Genre, dtype: object

```

Figure 15: Cleaning up the genre column

After cleaning up the data, I begin to analyze the dataset. First I determine what are the top five most common genre movies that made it to this list. I use ``value_counts().head(5)`` to display the top 5 results. Figure 16 shows the code and results. The top movie genres are drama, comedy, romance, thriller, and crime.



Figure 16: Determining top movie genres

I then use ``sort_values`` to sort the DataFrame by the Number of votes column to determine what top 5 movies had the most votes. Figure 17 shows the code and results.

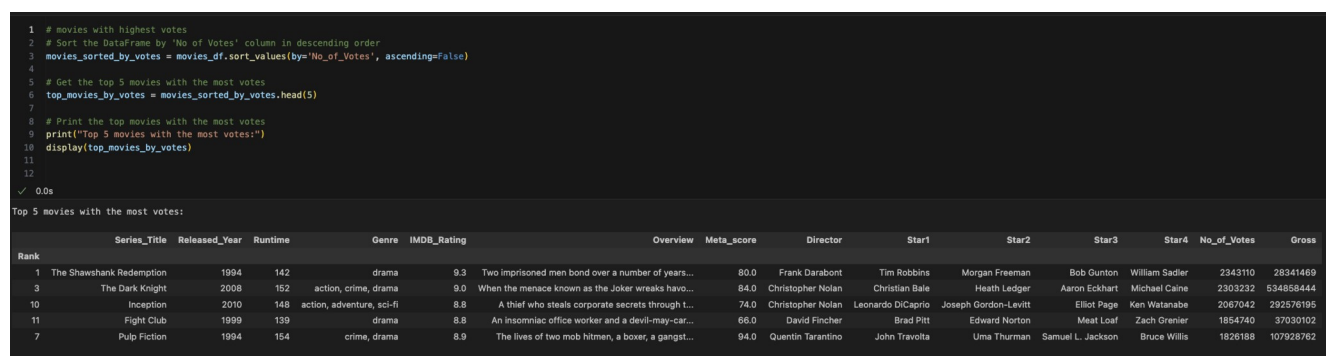


Figure 17: Determining top 5 most voted movies

I then do the same thing but to see which movies had the least number of votes. Figure 18 shows the results. The least number of votes is 25,198. This lets me know that none of the movies have a small number of votes that could truly skew the data such as one 5 star vote that makes the movie rank in the top ten. However, this does not mean the results aren't skewed. Further analysis would need to be completed to correctly determine this.

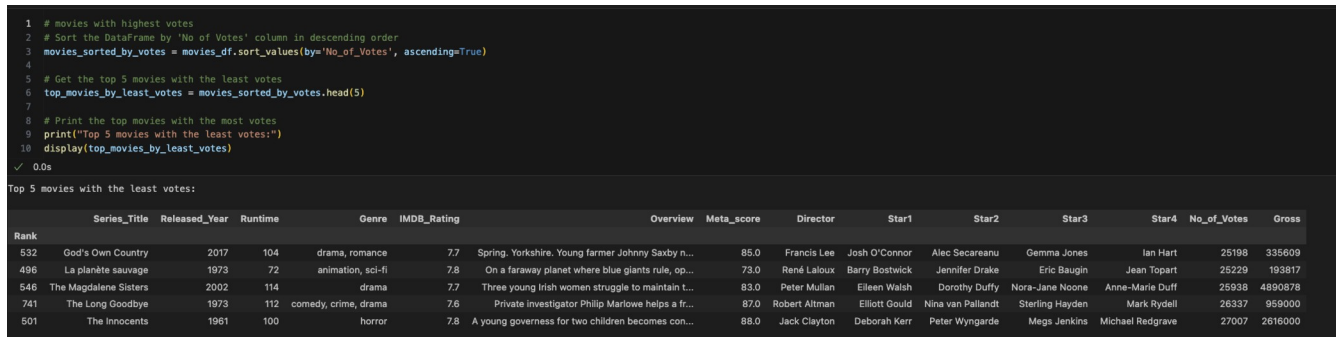


Figure 18: Determining bottom 5 least voted movies

The next question I answered was which movies had the longest runtime and which movies had the shortest runtime. I used the built-in `nlargest()` and `nsmallest()` to determine this. Figures 19 and 20 show the respective code along with the results.

```

1 #which movies have the longest runtime
2 # Find the movie with the longest runtime
3 movie_longest_runtime = movies_df.nlargest(1, 'Runtime')
4
5 # Print the movie details
6 print("Movie with the longest runtime:")
7 movie_longest_runtime
8

```

✓ 0.0s

Movie with the longest runtime:

	Series_Title	Released_Year	Runtime	Genre	IMDB_Ra
Rank					
209	Gone with the Wind	1939	238	drama, history, romance	

Figure 19: Determining which movie has the longest runtime

```

1 #movie with shortest runtime
2 # Find the movie with the shortest runtime
3 movie_shortest_runtime = movies_df.nsmallest(1, 'Runtime')
4
5 # Print the movie details
6 print("Movie with the shortest runtime:")
7 movie_shortest_runtime
8
✓ 0.0s

```

Movie with the shortest runtime:

Rank	Series_Title	Released_Year	Runtime	Genre	IMDB_Rating
680	The Secret of Kells	2009	71	animation, adventure, family	

Figure 20: Determining which movie has the shortest runtime

To demonstrate the usefulness of `query()`, I made a very simple example that filters the movie dataset. I have the code print out the movies that have a released year past 2004, with an IMDB rating greater than 8.0 and a metascore higher than 90.0. For this only 11 movies meet the criteria, and three of the eleven are animated movies. Figure 21 shows the results.

```

1 # showing query
2 filtered_movies = movies_df.query('IMDB_Rating > 8.0 and Released_Year > 2004 and Meta_score > 90.0')
3 filtered_movies

```

✓ 0.0s

Rank	Series_Title	Released_Year	Runtime	Genre	IMDB_Rating	Overview	Meta_score	Director
21	Gisaengchung	2019	132	comedy, drama, thriller	8.6	Greed and class discrimination threaten the ne...	96.0	Bong Jo
57	WALL-E	2008	98	animation, adventure, family	8.4	In the distant future, a small waste-collectin...	95.0	Andrew S
72	Jodaeiye Nader az Simin	2011	123	drama	8.3	A married couple are faced with a difficult de...	95.0	Asghar F
114	Toy Story 3	2010	103	animation, adventure, comedy	8.2	The toys are mistakenly delivered to a day-car...	92.0	Lee U
115	Pan's Labyrinth	2006	118	drama, fantasy, war	8.2	In the Falangist Spain of 1944, the bookish yo...	98.0	Guillermo de
116	There Will Be Blood	2007	158	drama	8.2	A story of family, religion, hatred, oil and m...	93.0	Paul Thomas And
143	Spotlight	2015	129	biography, crime, drama	8.1	The true story of how the Boston Globe uncover...	93.0	Tom Mc
145	Portrait de la jeune fille en feu	2019	122	drama, romance	8.1	On an isolated island in Brittany at the end o...	95.0	Céline Sc
146	12 Years a Slave	2013	134	biography, drama, history	8.1	In the antebellum United States, Solomon North...	96.0	Steve Mc
153	Inside Out	2015	95	animation, adventure, comedy	8.1	After young Riley is uprooted from her Midwest...	94.0	Pete
163	No Country for Old Men	2007	122	crime, drama, thriller	8.1	Violence and mayhem ensue after a hunter stumb...	91.0	Ethan

Figure 21: Using `query()` to filter data

To demonstrate using `eval`, I calculated what the total earnings would be by multiplying the gross column by the number of votes. This is not going to be realistic or is necessary. This is simply to show the usefulness of using `eval()`. Figure 22 shows the code and results.

```
1 # showing eval()
2 df = movies_df.eval('Total_Earnings = Gross * No_of_Votes')
3 df.head(5)
```

✓ 0.0s

Rank	Series_Title	Released_Year	Runtime	Genre	IMDB_Rating	Overview	Meta_score	Director	Star1	Star2	Star3	Star4	No_of_Votes	Gross	Total_Earnings
1	The Shawshank Redemption	1994	142	drama	9.3	Two imprisoned men bond over a number of years...	80.0	Frank Darabont	Tim Robbins	Morgan Freeman	Bob Gunton	William Sadler	2343110	28341469	66407179428590
2	The Godfather	1972	175	crime, drama	9.2	An organized crime dynasty's aging patriarch...	100.0	Francis Ford Coppola	Marlon Brando	Al Pacino	James Caan	Diane Keaton	1620367	134966411	218695118492837
3	The Dark Knight	2008	152	action, crime, drama	9.0	When the menace known as the Joker wreaks havoc...	84.0	Christopher Nolan	Christian Bale	Heath Ledger	Aaron Eckhart	Michael Caine	2303232	534858444	1231903083691008
4	The Godfather: Part II	1974	202	crime, drama	9.0	The early life and career of Vito Corleone in ...	90.0	Francis Ford Coppola	Al Pacino	Robert De Niro	Robert Duvall	Diane Keaton	1129952	57300000	64746249600000
5	12 Angry Men	1957	96	crime, drama	9.0	A jury holdout attempts to prevent a miscarria...	96.0	Sidney Lumet	Henry Fonda	Lee J. Cobb	Martin Balsam	John Fiedler	689845	4360000	3007724200000

Figure 22: Using `eval()` to calculate new data

The next question I answer is which year had the most top movies released. I accomplish this by using the `groupby()` method and `count()` method. I store the results in a variable named `best_year`. From this information I determined that 2014 had the most released movies that were considered top movies. Figure 23 shows this code and output.

```
1 # determine which year had the most released movies
2 best_year = movies_df.groupby('Released_Year').count()['Series_Title'].nlargest(10)
3 #best_year.shape
4 pd.set_option('display.max_rows', None)
5 display(best_year)
6 pd.reset_option('display.max_rows')
```

✓ 0.0s

Released_Year	
2014	29
2004	28
2001	24
2009	24
2013	24
2007	23
2006	22
2003	21
1993	20
2010	20

Name: Series_Title, dtype: int64

Figure 23: Determining which year had the most top movies

Next, I determine which movie was the highest grossing movie. To accomplish this task, I used `idxmax()` to get the index value of the movie that has the largest number in the column `Gross`. Funnily

enough the highest grossing movie is ranked 332 with an IMDB rating of 7.9. Figure 24 shows the results.



```
1 # which movies had the highest gross
2 max_gross_index = movies_df['Gross'].idxmax() #get index of movie with highest gross
3 movie_with_highest_gross = movies_df.loc[max_gross_index, 'Series_Title']
4 highest_gross = movies_df.loc[max_gross_index, 'Gross']
5 rating = movies_df.loc[max_gross_index, 'IMDB_Rating']
6 print(f"The highest grossing movie was: {movie_with_highest_gross} \n"
7       f"bringing in ${highest_gross} and is ranked number {max_gross_index}"
8       f"in the top movies with an IMDb rating of {rating}")
```

[92] ✓ 0.0s

... The highest grossing movie was: Star Wars: Episode VII – The Force Awakens
bringing in \$936662225 and is ranked number 332 in the top movies with an IMDb rating of 7.9

Figure 24: Determining highest grossing movie

I next determine which movie genres have the highest ratings. I grouped the movies by genre and then used `agg()` to look at the mean of the IMDB ratings and for fun the sum of how much money the genre brings in. Interestingly, the category of *crime, mystery, thriller* has the highest rating at 8.50 but *action, adventure, sci-fi* brings in the most money. Figure 25 shows the results.


```

1 # Grouping the movies and looking at
2 genre_group = movies_df.groupby('Genre').agg({'IMDB_Rating': 'mean', 'Gross': 'sum'})
3 print(genre_group.sort_values(by='IMDB_Rating', ascending=False).head(10))
4 print("\n")
5 print(genre_group.sort_values(by='Gross', ascending=False).head(10))

```

✓ 0.0s

Genre	IMDB_Rating	Gross
crime, mystery, thriller	8.50	23341568
action, sci-fi	8.40	414723280
horror, sci-fi	8.40	78900000
drama, mystery, war	8.35	90328607
western	8.35	58221508
mystery, romance, thriller	8.30	3200000
crime, drama, sci-fi	8.30	6207725
comedy, musical, romance	8.30	8819028
adventure, mystery, thriller	8.30	13275000
crime, drama, fantasy	8.25	138923439

Genre	IMDB_Rating	Gross
action, adventure, sci-fi	7.928571	5898659459
animation, adventure, comedy	7.925000	4503337598
action, adventure, drama	8.225000	2668835626
action, adventure, fantasy	8.200000	2116341031
drama	7.942424	2010111145
animation, action, adventure	7.960000	2008694522
drama, romance	7.929630	1962226898
action, adventure, comedy	7.828571	1911202048
action, crime, drama	7.845000	1308133687
crime, drama, thriller	7.933333	1255884472

Figure 25: Determining which genre has the IMDB rating and which has the highest gross

I created a pivot table to analyze the movie ratings by director to help identify which directors consistently produce highly rated movies. I filtered the data to ensure only directors who have made four or more movies would be included. Figure 26 shows the code and the results.

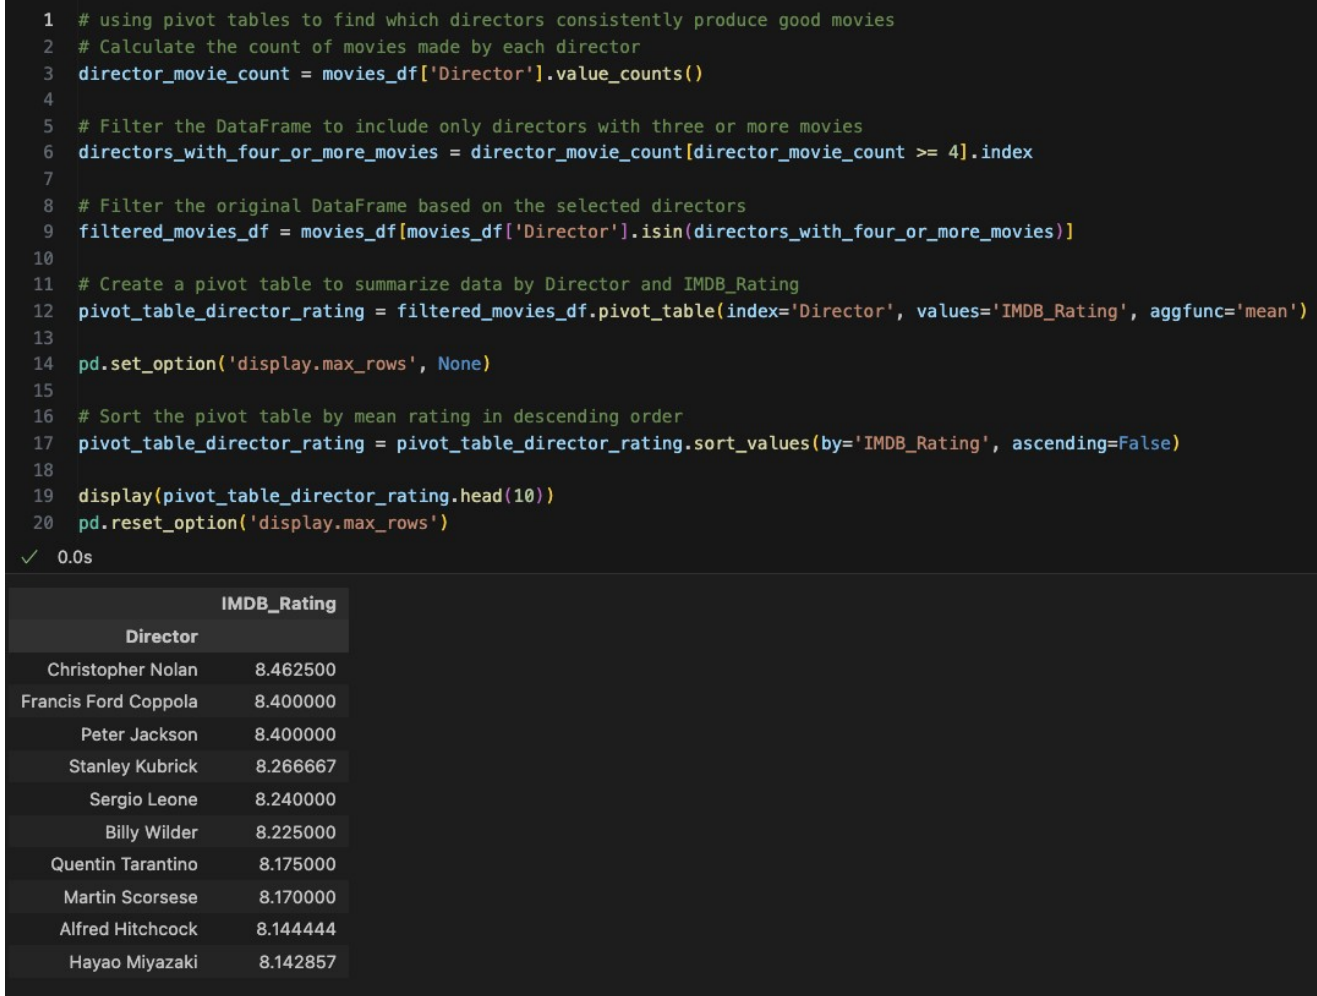


Figure 26: Discovering which directors consistently produce top movies using pivot tables

To demonstrate hierarchical indexing, I created a multi-index based on the genre and release year of the movies. This was accomplished by setting the index to 'Genre' and 'Released_Year'. Then an example was written to access the data within the DataFrame using hierarchical indexing. Figure 27 shows the code and the output.

```

1 # Set MultiIndex based on Genre and Release Year
2 movies_df.set_index(['Genre', 'Released_Year'], inplace=True)
3
4 # Sort the index
5 movies_df.sort_index(inplace=True)
6
7 # Display the DataFrame with MultiIndex
8 display(movies_df.head(20))
9
10 # Accessing data using MultiIndex
11 action_movies_2010 = movies_df.loc(['action, adventure, comedy', 2017])
12 display(action_movies_2010)
13

```

✓ 0.0s

Genre	Released_Year	Series_Title	Runtime	IMDB_Rating	Overview	Meta_score
action, adventure	1981	Raiders of the Lost Ark	115	8.4	In 1936, archaeologist and adventurer Indiana ...	85.0
	1982	First Blood	93	7.7	A veteran Green Beret is forced by a cruel She...	61.0
	1989	Indiana Jones and the Last Crusade	127	8.2	In 1938, after his father Professor Henry Jone...	65.0
	2005	Batman Begins	140	8.2	After training with his mentor, Batman begins ...	70.0
	2012	The Dark Knight Rises	164	8.4	Eight years after the Joker's reign of anarchy...	78.0
action, adventure, comedy	1980	The Blues Brothers	133	7.9	Jake Blues, just released from prison, puts to...	60.0
	2014	Guardians of the Galaxy	121	8.0	A group of intergalactic criminals must pull t...	76.0
	2014	Kingsman: The Secret Service	129	7.7	A spy organisation recruits a promising street...	60.0
	2016	Deadpool	108	8.0	A wisecracking mercenary gets experimented on ...	65.0
	2017	Thor: Ragnarok	130	7.9	Imprisoned on the planet Sakaar, Thor must rac...	74.0
	2017	Guardians of the Galaxy Vol. 2	136	7.6	The Guardians struggle to keep together as a t...	67.0
	2018	Deadpool 2	119	7.7	Foul-mouthed mutant mercenary Wade Wilson (a.k...	66.0
action, adventure, drama	1954	Shichinin no samurai	207	8.6	A poor village under attack by bandits recruit...	98.0
	1992	The Last of the Mohicans	112	7.7	Three trappers protect the daughters of a Brit...	76.0
	2000	Gladiator	155	8.5	A former Roman General sets out to exact venge...	67.0
	2001	The Lord of the Rings: The Fellowship of the Ring	178	8.8	A meek Hobbit from the Shire and eight compani...	92.0
	2002	The Lord of the Rings: The Two Towers	179	8.7	While Frodo and Sam edge closer to Mordor with...	87.0
	2002	The Count of Monte Cristo	131	7.8	A young man, falsely imprisoned by his jealous...	61.0
	2003	The Lord of the Rings: The Return of the King	201	8.9	Gandalf and Aragorn lead the World of Men agai...	94.0
	2006	Letters from Iwo Jima	141	7.9	The story of the battle of Iwo Jima between th...	89.0

Genre	Released_Year	Series_Title	Runtime	IMDB_Rating	Overview	Meta_score	Director
action, adventure, comedy	2017	Thor: Ragnarok	130	7.9	Imprisoned on the planet Sakaar, Thor must rac...	74.0	Taika Waititi
	2017	Guardians of the Galaxy Vol. 2	136	7.6	The Guardians struggle to keep together as a t...	67.0	James Gunn

Figure 27: Demonstrating hierarchical indexing

In order to demonstrate merging files, two new data sets were imported, TMDB 5000 Credits.csv and TMDB 5000 Movies.csv(see Figure 28-29). The datasets were then merged on the movie id and the new DataFrame was created(see Figure 30). Two columns were removed from the newly merged data(see Figure 31) and the index was set to the genre and release date column(see Figure 32). To demonstrate vectorized string operations, the first word of all the movie titles was extracted and added to a new column(see figure 33-34).

Merging Files example

```
1 # loading new datasets
2 df1 = pd.read_csv("tmdb_5000_credits.csv")
3 df2 = pd.read_csv("tmdb_5000_movies.csv")
4 df1.info()
```

✓ 0.2s

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   movie_id    4803 non-null   int64
1   title       4803 non-null   object
2   cast        4803 non-null   object
3   crew        4803 non-null   object
dtypes: int64(1), object(3)
memory usage: 150.2+ KB
```

Figure 28: Importing two new datasets and showing information about dataset 1

```
1 df2.info()
```

✓ 0.0s

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 20 columns):
#   Column              Non-Null Count  Dtype
---  -
0   budget              4803 non-null   int64
1   genres              4803 non-null   object
2   homepage            1712 non-null   object
3   id                  4803 non-null   int64
4   keywords            4803 non-null   object
5   original_language   4803 non-null   object
6   original_title      4803 non-null   object
7   overview            4800 non-null   object
8   popularity          4803 non-null   float64
9   production_companies 4803 non-null   object
10  production_countries 4803 non-null   object
11  release_date        4802 non-null   object
12  revenue             4803 non-null   int64
13  runtime             4801 non-null   float64
14  spoken_languages    4803 non-null   object
15  status              4803 non-null   object
16  tagline              3959 non-null   object
17  title               4803 non-null   object
18  vote_average        4803 non-null   float64
19  vote_count          4803 non-null   int64
dtypes: float64(3), int64(4), object(13)
memory usage: 750.6+ KB
```

Figure 29: Dataset 2 information

```

1 df1.columns = ['id', 'title', 'cast', 'crew']
2 merged_df = df2.merge(df1, on='id')
3 merged_df.info()

```

✓ 0.0s

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4803 entries, 0 to 4802

Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
0	budget	4803 non-null	int64
1	genres	4803 non-null	object
2	homepage	1712 non-null	object
3	id	4803 non-null	int64
4	keywords	4803 non-null	object
5	original_language	4803 non-null	object
6	original_title	4803 non-null	object
7	overview	4800 non-null	object
8	popularity	4803 non-null	float64
9	production_companies	4803 non-null	object
10	production_countries	4803 non-null	object
11	release_date	4802 non-null	object
12	revenue	4803 non-null	int64
13	runtime	4801 non-null	float64
14	spoken_languages	4803 non-null	object
15	status	4803 non-null	object
16	tagline	3959 non-null	object
17	title_x	4803 non-null	object
18	vote_average	4803 non-null	float64
19	vote_count	4803 non-null	int64
...			
21	cast	4803 non-null	object
22	crew	4803 non-null	object

dtypes: float64(3), int64(4), object(16)

memory usage: 863.2+ KB

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust

Figure 30: Merging the two datasets on 'id'

```

1 merged_df.drop(['title_y', 'homepage'], axis=1, inplace=True)
2 merged_df.info()

```

✓ 0.0s

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   budget                               4803 non-null   int64
1   genres                               4803 non-null   object
2   id                                    4803 non-null   int64
3   keywords                             4803 non-null   object
4   original_language                    4803 non-null   object
5   original_title                       4803 non-null   object
6   overview                             4800 non-null   object
7   popularity                           4803 non-null   float64
8   production_companies                 4803 non-null   object
9   production_countries                 4803 non-null   object
10  release_date                         4802 non-null   object
11  revenue                              4803 non-null   int64
12  runtime                             4801 non-null   float64
13  spoken_languages                     4803 non-null   object
14  status                               4803 non-null   object
15  tagline                              3959 non-null   object
16  title_x                              4803 non-null   object
17  vote_average                         4803 non-null   float64
18  vote_count                           4803 non-null   int64
19  cast                                 4803 non-null   object
20  crew                                 4803 non-null   object
dtypes: float64(3), int64(4), object(14)
memory usage: 788.1+ KB

```

Figure 31: Removing unused columns


```

1 merged_df.set_index(['genres', 'release_date'], inplace=True)
2 merged_df.head()

```

✓ 0.0s

genres	release_date	budget	id	keywords	original_language	original_title	overview	popularity
[{"id": 28, "name": "Action"}, {"id": 12, "name": "Adventure"}, {"id": 14, "name": "Fantasy"}, {"id": 878, "name": "Science Fiction"}]	2009-12-10	237000000	19995	[{"id": 1463, "name": "culture clash"}, {"id": ...}]	en	Avatar	In the 22nd century, a paraplegic Marine is di...	150.437577
[{"id": 12, "name": "Adventure"}, {"id": 14, "name": "Fantasy"}, {"id": 28, "name": "Action"}]	2007-05-19	300000000	285	[{"id": 270, "name": "ocean"}, {"id": 726, "name": "na..."}]	en	Pirates of the Caribbean: At World's End	Captain Barbossa, long believed to be dead, ha...	139.082615
[{"id": 28, "name": "Action"}, {"id": 12, "name": "Adventure"}, {"id": 80, "name": "Crime"}]	2015-10-26	245000000	206647	[{"id": 470, "name": "spy"}, {"id": 818, "name": "name..."}]	en	Spectre	A cryptic message from Bond's past sends him o...	107.376788
[{"id": 28, "name": "Action"}, {"id": 80, "name": "Crime"}, {"id": 18, "name": "Drama"}, {"id": 52, "name": "Mystery"}]	2012-07-16	250000000	49026	[{"id": 849, "name": "dc comics"}, {"id": 853, "name": "na..."}]	en	The Dark Knight Rises	Following the death of District Attorney Harve...	112.312950

Figure 32: Setting the index to 'genres' and 'release_date'

```

1 # vectorized string operations
2 merged_df['first_word'] = merged_df['title_x'].str.split().str[0]
3 merged_df

```

✓ 0.0s

Figure 33: Code for vectorized string operations

ogline	title_x	vote_average	vote_count	cast	crew	first_word
er the world of ndora.	Avatar	7.2	11800	[{"cast_id": 242, "character": "Jake Sully", "...	[{"credit_id": "52fe48009251416c750aca23", "de...	Avatar
e end of the d, the venture egins.	Pirates of the Caribbean: At World's End	6.9	4500	[{"cast_id": 4, "character": "Captain Jack Spa...	[{"credit_id": "52fe4232c3a36847f800b579", "de...	Pirates
an No One capes	Spectre	6.3	4466	[{"cast_id": 1, "character": "James Bond", "cr...	[{"credit_id": "54805967c3a36829b5002c41", "de...	Spectre
The Legend Ends	The Dark Knight Rises	7.6	9106	[{"cast_id": 2, "character": "Bruce Wayne / Ba...	[{"credit_id": "52fe4781c3a36847f81398c3", "de...	The
in our world, und in other.	John Carter	6.1	2124	[{"cast_id": 5, "character": "John Carter", "c...	[{"credit_id": "52fe479ac3a36847f813eaa3", "de...	John

Figure 34: Output to Figure 32 code

To demonstrate the `concat()` method, I split the `movie_df` dataset into two halves stored in two different variables. I then combined the two halves back together in a new DataFrame using `concat()`.

To verify this worked, I print the shape of the new DataFrame to ensure it is the same shape as the original DataFrame. Figure 35 shows the code and output.

```
1 # Example using concat to combine DataFrames along rows
2 movies_part1 = movies_df.iloc[:500] # First half of the DataFrame
3 movies_part2 = movies_df.iloc[500:] # Second half of the DataFrame
4
5 # Concatenate the two parts
6 movies_combined = pd.concat([movies_part1, movies_part2])
7
8 # Verify the concatenation
9 print(movies_combined.shape)
10
```

✓ 0.0s

(749, 12)

Figure 35: Demonstrating `concat()`

To demonstrate using `concat()` to append new rows to an existing DataFrame, I define two generic Movies to add to the end of the list. Figure 36 shows the code and the output after appending the two new movies to the original dataset of `movies_df`.

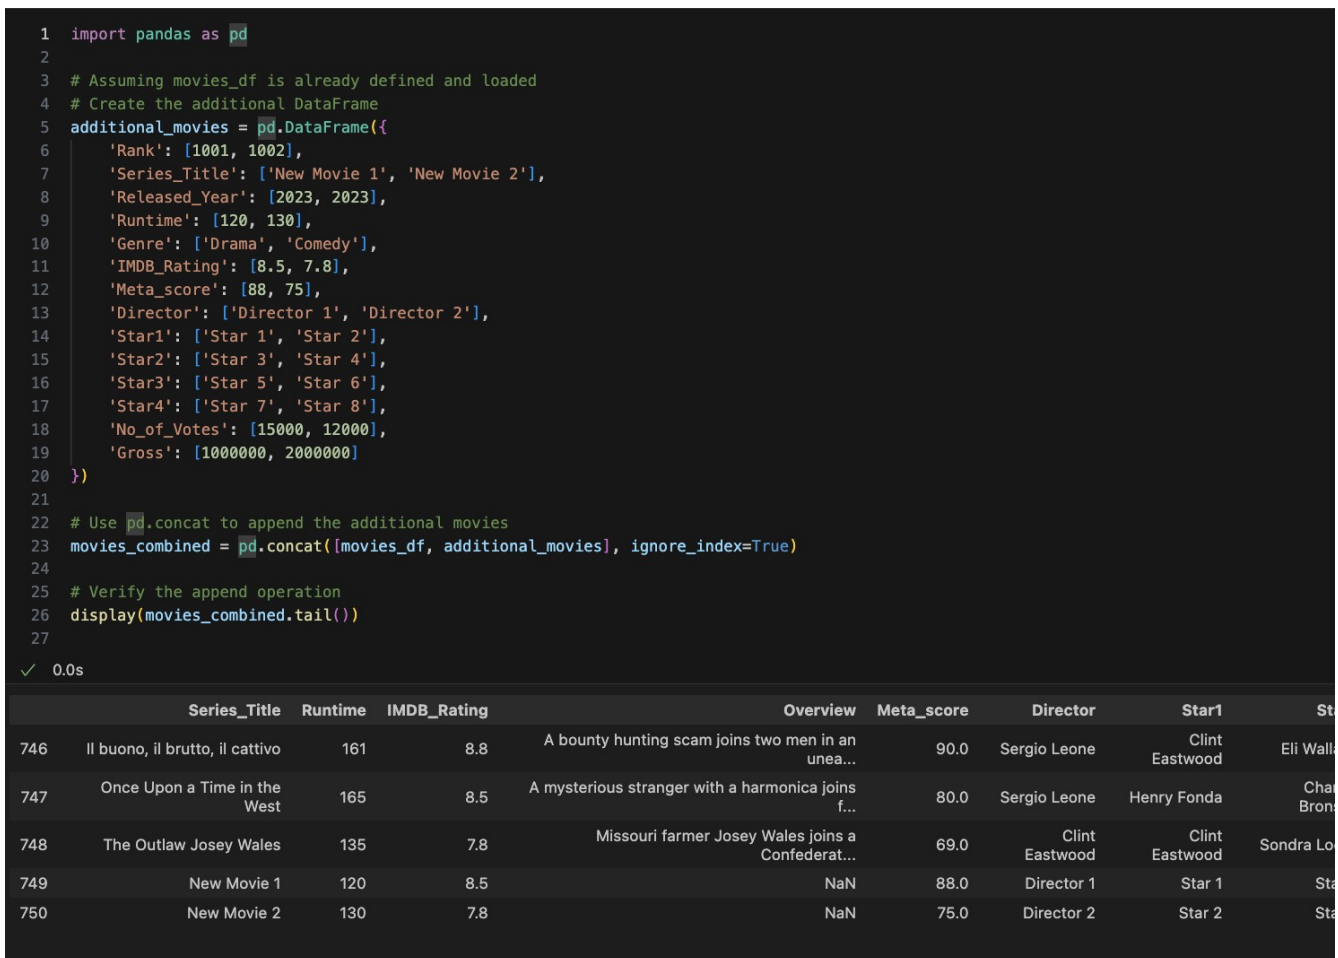


Figure 36: Using `concat()` to append new data

3. Conclusion

This report documents my journey in learning Pandas, a powerful data manipulation library in Python. Key concepts I explored include data structures essential for handling and analyzing structured data. I learned to perform various data operations, including data indexing, merging datasets, grouping data and more.

By using real data, I was able to explore how to go about cleaning the data up properly before beginning to analyze it. Many errors I encountered were related to missing values and data types, which are significantly different from syntax errors I encountered in previous projects. However, through practice and persistence, I was able to clean the data up and obtain datasets that could be analyzed properly using Pandas.

References

1. Harshit Shankhdhar. (2021). IMDB Dataset of Top 1000 Movies and TV Shows. Retrieved from <https://www.kaggle.com/datasets/harshitshankhdhar/IMDB-dataset-of-top-1000-movies-and-tv-shows>
2. The Movie Database (TMDB). (2018). TMDB 5000 Movie Dataset. Retrieved from <https://www.kaggle.com/datasets/tmdb/tmdb-movie-metadata>
3. VanderPlas, J. (2016). *Python Data Science Handbook*. O'Reilly Media. Retrieved from <https://jakevdp.github.io/PythonDataScienceHandbook/index.html>